Rethinking Kubernetes:

VS

VPA

PerfectScale

The main weakness of VPA is:

1. Need to maintain and tune CRDs. for each and every workload together with Prometheus.

2. VPA does not work with HPA

3. The algorithm can lead to suboptimal results (historical data limitation etc)

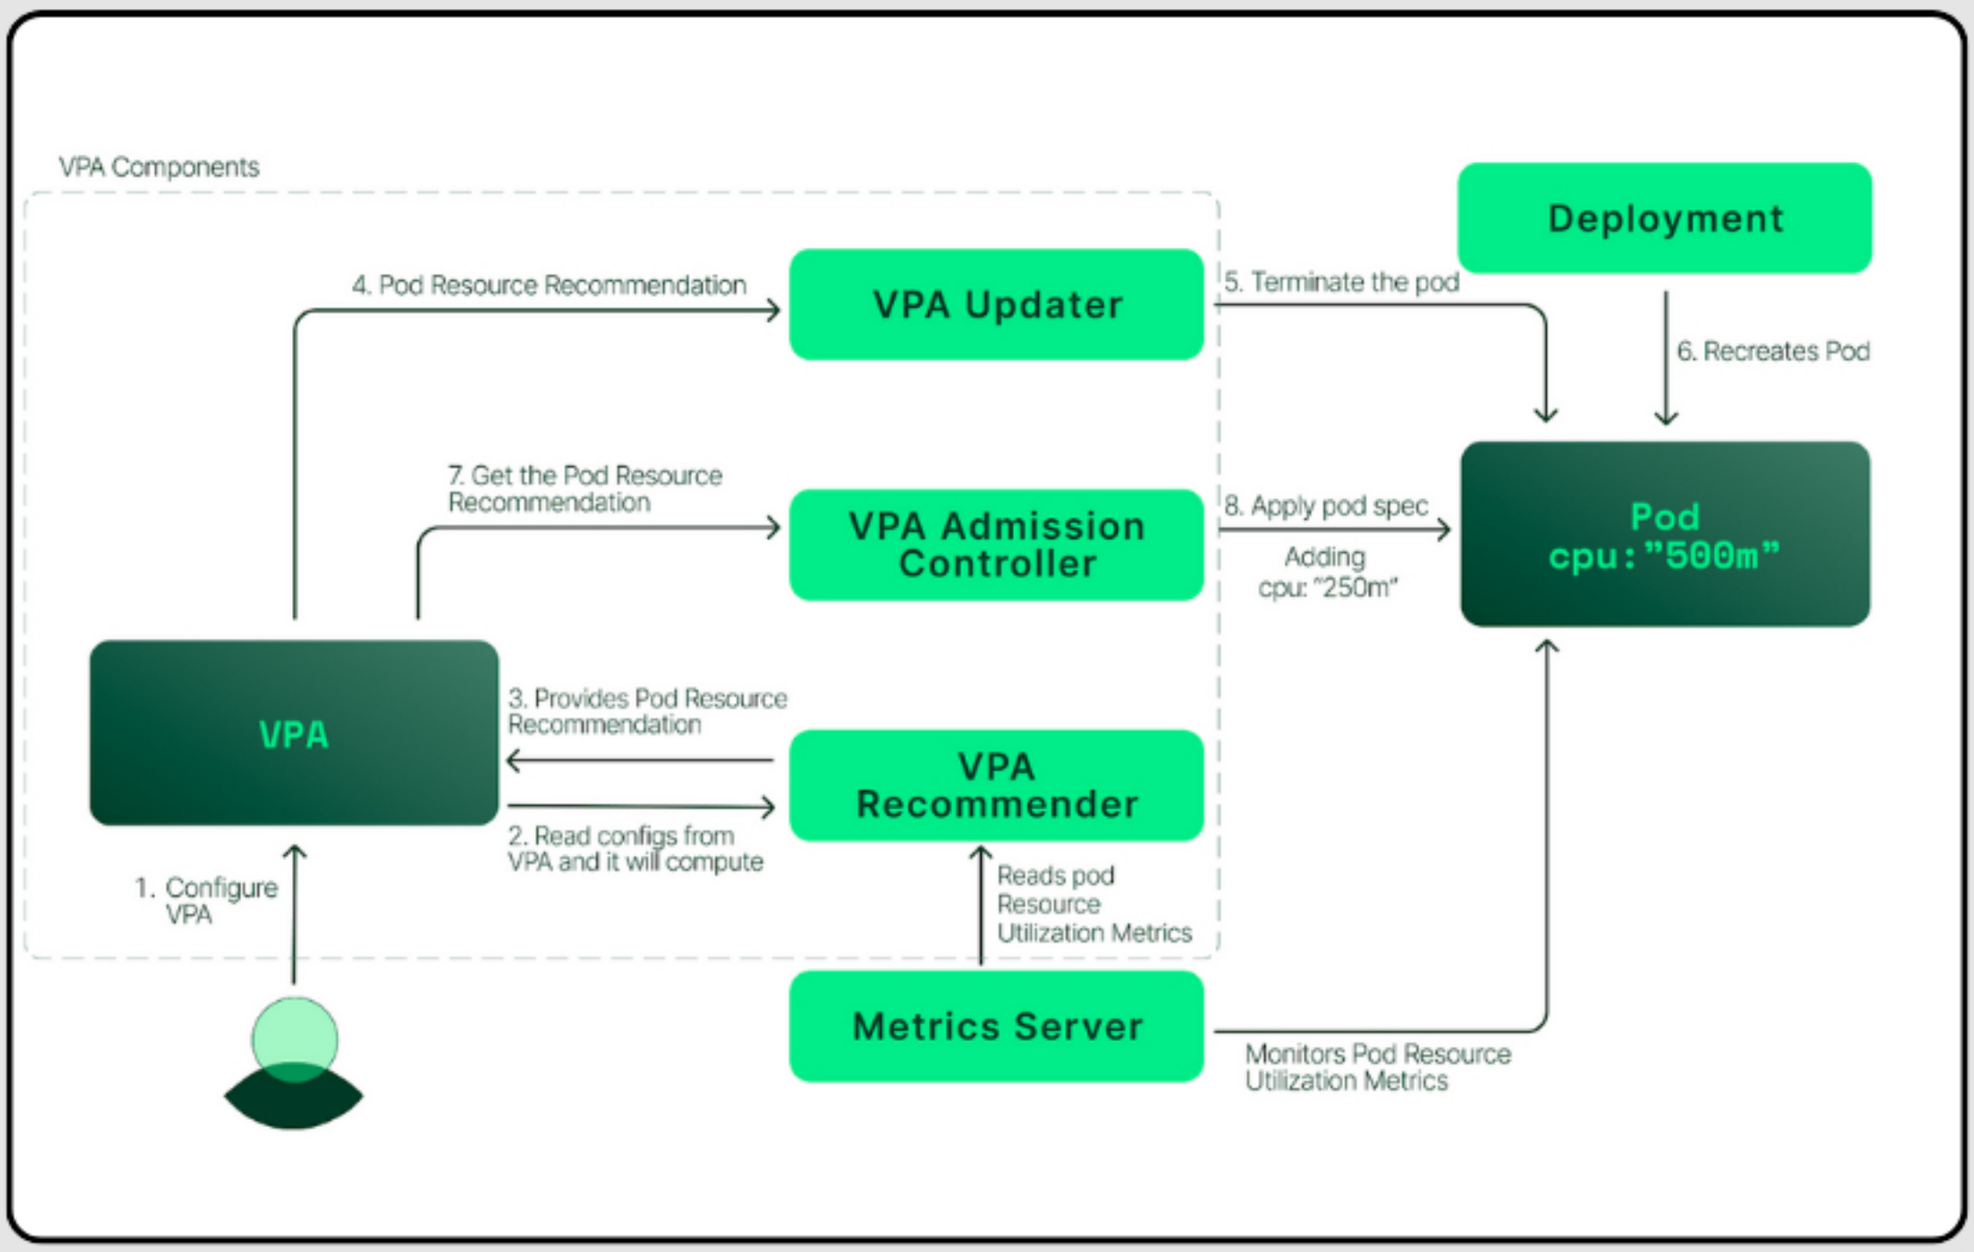

PerfectScale

Typical adoption path with the Vertical Pod Autoscaler (VPA):

PerfectScale

Day 0 [Preparation]

Set up Prometheus and Metrics Servers.
Prepare a CRD for the first service with very precise VPA parameters.

PerfectScale

Day 1 [Implementation]

Repeat the same process for dozens or hundreds of workloads across multiple clusters (VPA parameters are usually not detailed).

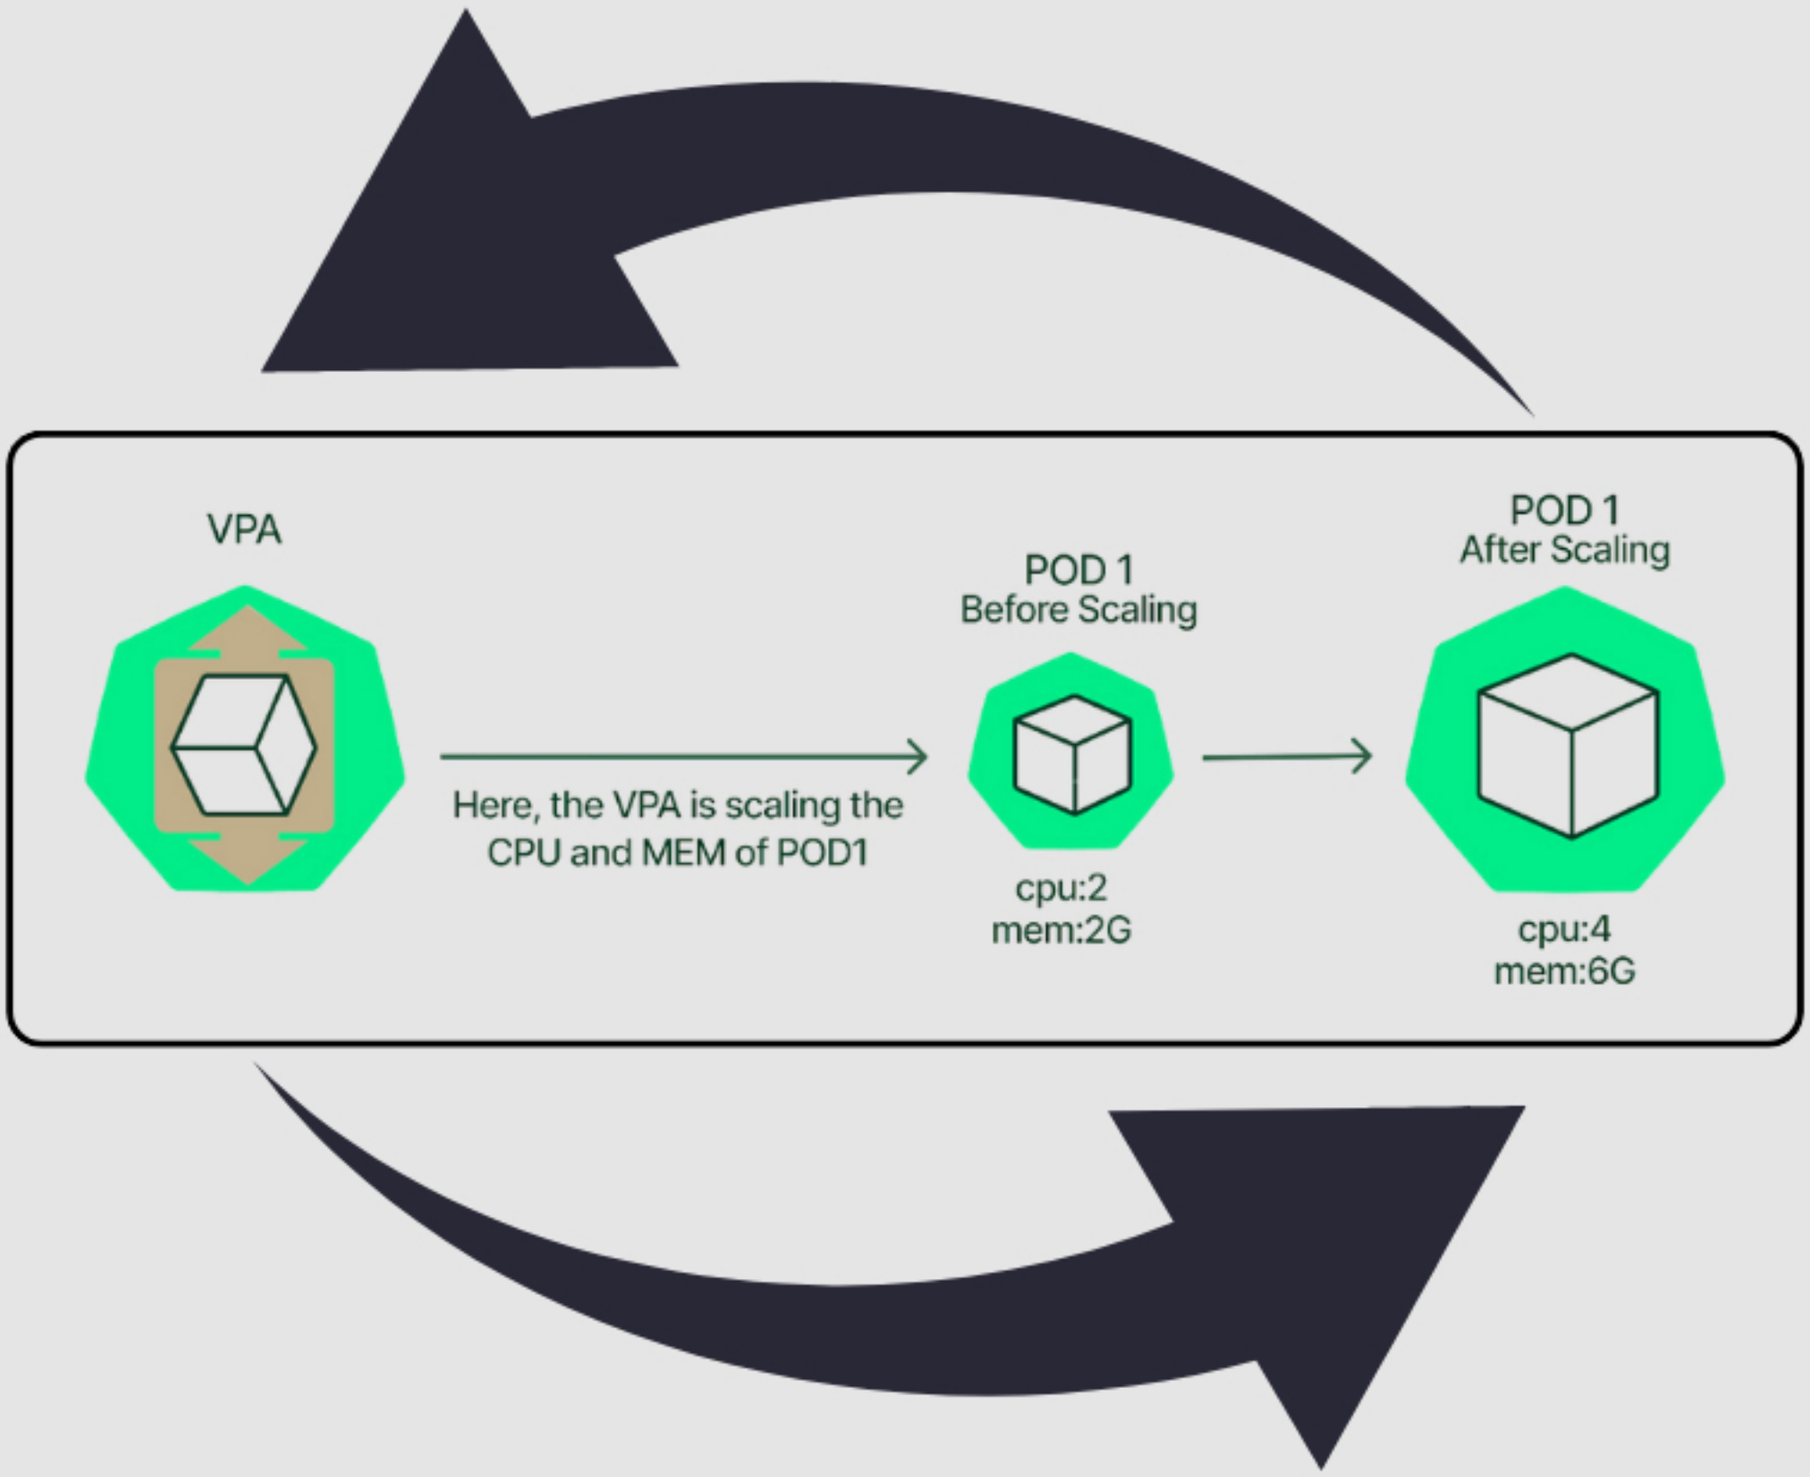

PerfectScale

Day 2 [Operation]

Continuously add CRD to any new workload, providing precise VPA parameters.

Pause VPA every time you deploy code with significant changes (e.g., adding internal cache, introducing heavy functions).

Manually fine-tune VPA parameters for each workload.

PerfectScale

Now, let's contrast this with the adoption path with PerfectScale (CRD moved under the hood):

PerfectScale

PerfectScale

Day 0 [Preparation]

Set up PerfectScale (2 min).

Day 1 [Implementation]

- Add a label to the workload to initiate automation. (including deployments with HPA enabled).
- Label the entire namespace to include all existing workloads.
- New workloads in the "labeled" namespace will be automatically included unless explicitly labeled "no automation."

Day 2 [Operation]

Enjoy

PerfectScale provides an effortless way to vertically scale K8s clusters without needing to manually fine-tune VPA parameters for each workload.

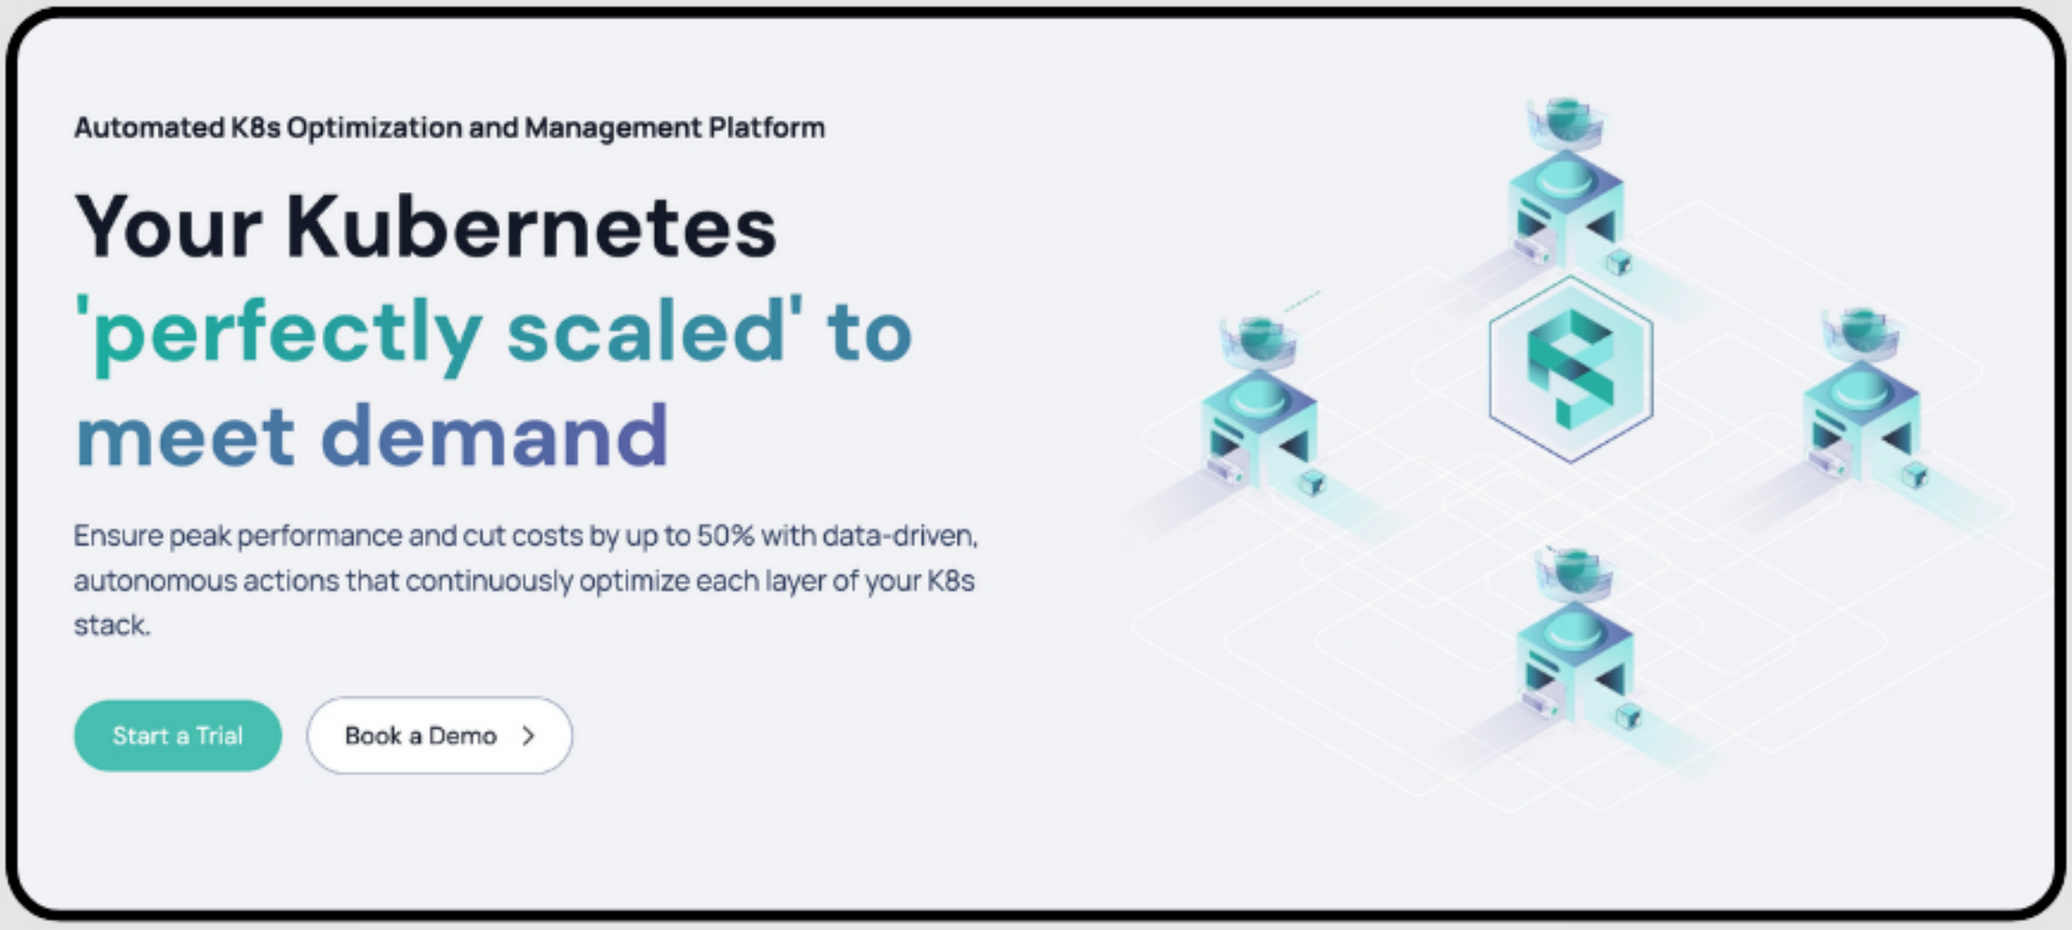

See how easy it is to autonomously right-size and scale your K8s environment with **PerfectScale.**

PerfectScale